

A dark, atmospheric photograph of the Golden Gate Bridge in San Francisco, viewed from a low angle looking down the length of the bridge towards the water. The bridge's iconic towers and suspension cables are visible against a hazy sky. The water below is dark, with a small sailboat visible in the distance.

DISASTER RELIEF

LOCATION

a GA capstone by naren rajasekaran

What is
DISASTER
RELIEF
LOCATION
about ?

disaster

dɪ'zɑːstə/

a sudden event resulting in great damage and loss of life

With the recent events around the world with regards to natural disasters (Hurricane Harvey, Irma) as well as terror attacks in Paris/London, people are increasingly turning to social media to report events and to request for help.

My project was to reduce the noise from one of these platforms (Twitter), provide tracking (approximate) and to assign locations for these relief workers.

Dataset

I used the Twitter API services to scrape relevant topics which were widely spoken of (Trending) and did my best to bring some form of localisation for it. This is to simulate the closeness it will have to disasters where it is usually based in a single location (region, state) and not the entire country. As such, I chose to collect tweets by reducing the geo-range to within the boundaries of the USA by providing the necessary bounding box coordinates.



DATE October 3, 2017

SUBJECT Re: Can't find coordinates in tweet

EMAIL
customersupport@twitter.com

Hi Customer Support Guy,
im using tweets for a datascience project but most tweets don't have geolocations. what gives?

DATE October 3, 2017

SUBJECT Re: Can't find coordinates in tweet

Hi Naren,
sucks to be you. lol

Problems with my dataset

A DRAMATIC
RETELLING OF
EVENTS

coordinates

Coordinates

Nullable. Represents the geographic location of this Tweet as reported by the user or client application.

```
"coordinates":  
{  
  "coordinates":  
    [  
      -75.14310264,  
      40.05701649  
    ],  
  "type": "Point"  
}
```


Coordinates are opt-in per tweet. Every tweet has to be configured to send current location while tweeting.

```
"place":  
{  
  "attributes": {},  
  "bounding_box":  
    {  
      "coordinates":  
        [[  
          [-77.119759, 38.791645],  
          [-76.909393, 38.791645],  
          [-76.909393, 38.995548],  
          [-77.119759, 38.995548]  
        ]],  
      "type": "Polygon"  
    },  
  "country": "United States",  
  "country_code": "US",  
  "full_name": "Washington, DC",  
  "id": "01fbe706f872cb32",  
  "name": "Washington",  
  "place_type": "city",  
  "url": "http://api.twitter.com/1/geo/id/01fbe706f872cb32.json"  
}
```

An approximate location is tied to the tweets by default. We are given bounding boxes and the state of the origin of the tweet.

This ended up being a blessing in disguise as in most rescue efforts especially involving people who have cellphone connectivity, they tend to turn off non-essential services including locations as they tend to drain the most battery.

Other issues encountered



API Calling and Cleaning

```
26 def tweet_search(api, query, max_tweets, max_id, since_id, geocode):
27     ''' Function that takes in a search string 'query', the maximum
28     number of tweets 'max_tweets', and the minimum (i.e., starting)
29     tweet id. It returns a list of tweepy.models.Status objects. '''
30
31     searched_tweets = []
32     while len(searched_tweets) < max_tweets:
33         remaining_tweets = max_tweets - len(searched_tweets)
34         try:
35             new_tweets = api.search(q=query, count=remaining_tweets,
36                                   since_id=str(since_id),
37                                   max_id=str(max_id-1))
38             # geocode=geocode
39             print('found', len(new_tweets), 'tweets')
40             if not new_tweets:
41                 print('no tweets found')
42                 break
43             searched_tweets.extend(new_tweets)
44             max_id = new_tweets[-1].id
45         except tweepy.TweepError:
46             print('exception raised, waiting 15 minutes')
47             print('until:', dt.datetime.now()+dt.timedelta(minutes=15), '')
48             time.sleep(15*60)
49             break # stop the loop
50     return searched_tweets, max_id
51
52
53 def get_tweet_id(api, date='', days_ago=9, query='a'):
54     ''' Function that gets the ID of a tweet. This ID can then be
55     used as a 'starting point' from which to search. The query is
56     required and has been set to a commonly used word by default.
57     The variable 'days_ago' has been initialized to the maximum
58     amount we are able to search back in time (9).'''
59
60     if date:
61         # return an ID from the start of the given day
62         td = date + dt.timedelta(days=1)
63         tweet_date = '{0}-{1:0>2}-{2:0>2}'.format(td.year, td.month, td.day)
64         tweet = api.search(q=query, count=1, until=tweet_date)
65     else:
66         # return an ID from __ days ago
67         td = dt.datetime.now() - dt.timedelta(days=days_ago)
68         tweet_date = '{0}-{1:0>2}-{2:0>2}'.format(td.year, td.month, td.day)
69         # get list of up to 10 tweets
```

my precious code to call Twitter

API



the reply i get



```
{
  "contributors": null,
  "truncated": false,
  "text": "RT @CarolKfanclub: No need for #280characters to describe this wonderful woman. A picture says it all... https://t.co/Hkonl5EXxA",
  "is_quote_status": false, "in_reply_to_status_id": null,
  "id": "928692404407894016", "favorite_count": 0,
  "entities": {
    "symbols": [], "user_mentions": [
      {
        "id": "525718207", "indices": [3, 17], "id_str": "525718207",
        "screen_name": "CarolKfanclub", "name": "CarolKirkwoodFanClub", "hashtags": [
          {
            "indices": [31, 45], "text": "280characters"
          }
        ], "urls": [],
        "media": [
          {
            "source_user_id": "525718207", "source_status_id_str": "928204281341202433", "expanded_url": "https://twitter.com/CarolKfanclub/status/928204281341202433/photo/1",
            "display_url": "pic.twitter.com/Hkonl5EXxA", "url": "https://t.co/Hkonl5EXxA", "media_url_https": "https://pbs.twimg.com/media/D0Gk3EYw4AAJFT4.jpg",
            "source_user_id_str": "525718207", "source_status_id": "928204281341202433", "id_str": "928204264299749376",
            "sizes": {
              "large": {
                "h": 973, "resize": "fit", "w": 720, "small": {
                  "h": 680, "resize": "fit", "w": 503,
                  "medium": {
                    "h": 973, "resize": "fit", "w": 720, "thumb": {
                      "h": 150, "resize": "crop", "w": 150
                    }
                  }, "indices": [105, 128], "type": "photo", "id": "928204264299749376", "media_url": "http://pbs.twimg.com"
                }
              }
            }, "source": "a href=http://www.twitter.com rel=nofollowa>Twitter for Windowsa</a>", "in_reply_to_screen_name": null, "in_reply_to_user_id": null,
            "retweet_count": 4, "id_str": "928692404407894016", "favorited": false, "retweeted_status": {
              "contributors": null,
              "truncated": false, "text": "No need for #280characters to describe this wonderful woman. A picture says it all... https://t.co/Hkonl5EXxA",
              "is_quote_status": false, "in_reply_to_status_id": null, "id": "928204281341202433", "favorite_count": 67, "entities": {
                "symbols": [],
                "user_mentions": [], "hashtags": [
                  {
                    "indices": [12, 26], "text": "280characters"
                  }
                ], "urls": [],
                "media": [
                  {
                    "source_user_id": "525718207", "source_status_id_str": "928204281341202433", "expanded_url": "https://twitter.com/CarolKfanclub/status/928204281341202433/photo/1",
                    "display_url": "pic.twitter.com/Hkonl5EXxA", "url": "https://t.co/Hkonl5EXxA", "media_url_https": "https://pbs.twimg.com/media/D0Gk3EYw4AAJFT4.jpg",
                    "id_str": "928204264299749376", "sizes": {
                      "large": {
                        "h": 973, "resize": "fit", "w": 720,
                        "small": {
                          "h": 680, "resize": "fit", "w": 503,
                          "medium": {
                            "h": 973, "resize": "fit", "w": 720,
                            "thumb": {
                              "h": 150, "resize": "crop", "w": 150
                            }
                          }
                        }, "indices": [86, 109], "type": "photo",
                        "id": "928204264299749376", "media_url": "http://pbs.twimg.com/media/D0Gk3EYw4AAJFT4.jpg"
                      }
                    }, "source": "a href=http://twitter.com/download/android rel=nofollowa>Twitter for Androida</a>",
                        "retweeted": false, "coordinates": null, "source": "a href=http://twitter.com/download/android rel=nofollowa>Twitter for Androida</a>",
                        "in_reply_to_screen_name": null, "in_reply_to_user_id": null, "retweet_count": 4, "id_str": "928204281341202433", "favorited": false,
                        "user": {
                          "follow_request_sent": false, "has_extended_profile": false, "profile_use_background_image": true,
                          "default_profile_image": false, "id": "525718207", "profile_background_image_url_https": "https://pbs.twimg.com/profile_background_images/378800000167381390/ayb-kUNF.jpeg",
                          "verified": false, "translator_type": "none",
                          "profile_text_color": "3C3940", "profile_image_url_https": "https://pbs.twimg.com/profile_images/925482767747567617/m24XjXm_normal.jpg",
                          "profile_sidebar_fill_color": "95E8EC", "entities": {
                            "urls": {
                              "urls": [
                                {
                                  "url": "https://t.co/hjrmYf8pUd", "indices": [0, 23],
                                  "expanded_url": "http://www.bbc.co.uk/programmes/b006v5tb/presenters/carol-kirkwood", "display_url": "bbc.co.uk/programmes/b006v5tb/presenters/carol-kirkwood", "description": {
                                    "urls": [], "followers_count":

```

This person retweeted someone who retweeted someone who retweeted someone.

Preprocessing / NLTK

```
76
77     # output sentiment
78
79     print "Total tweets",len(lis)
80     print "Positive ",float(p/cout)*100,"% "
81     print "Negative ",float(n/cout)*100,"% "
82     print "Neutral ",float(net/len(lis))*100,"% "
83     #print lis
84     # determine if sentiment is positive, negative, or neutral
85
86     # output sentiment
87     #print sentiment
88
```

In [5]:

```
1 lis
0.0,
0.0,
0.14545454545454545,
0.5,
0.0,
0.5,
0.14545454545454545,
-0.8,
0.5,
0.14545454545454545,
0.0,
0.0,
0.175,
0.4681818181818182,
0.0,
0.0,
-0.1,
0.5,
0.0,
0.0,
```

Ignore positive tweets/retweets/focus on negative

EDA \\\

36000 TWEETS PER
HASHTAG

3%

WITH COORDINATES

Trimming - We are left with

- id

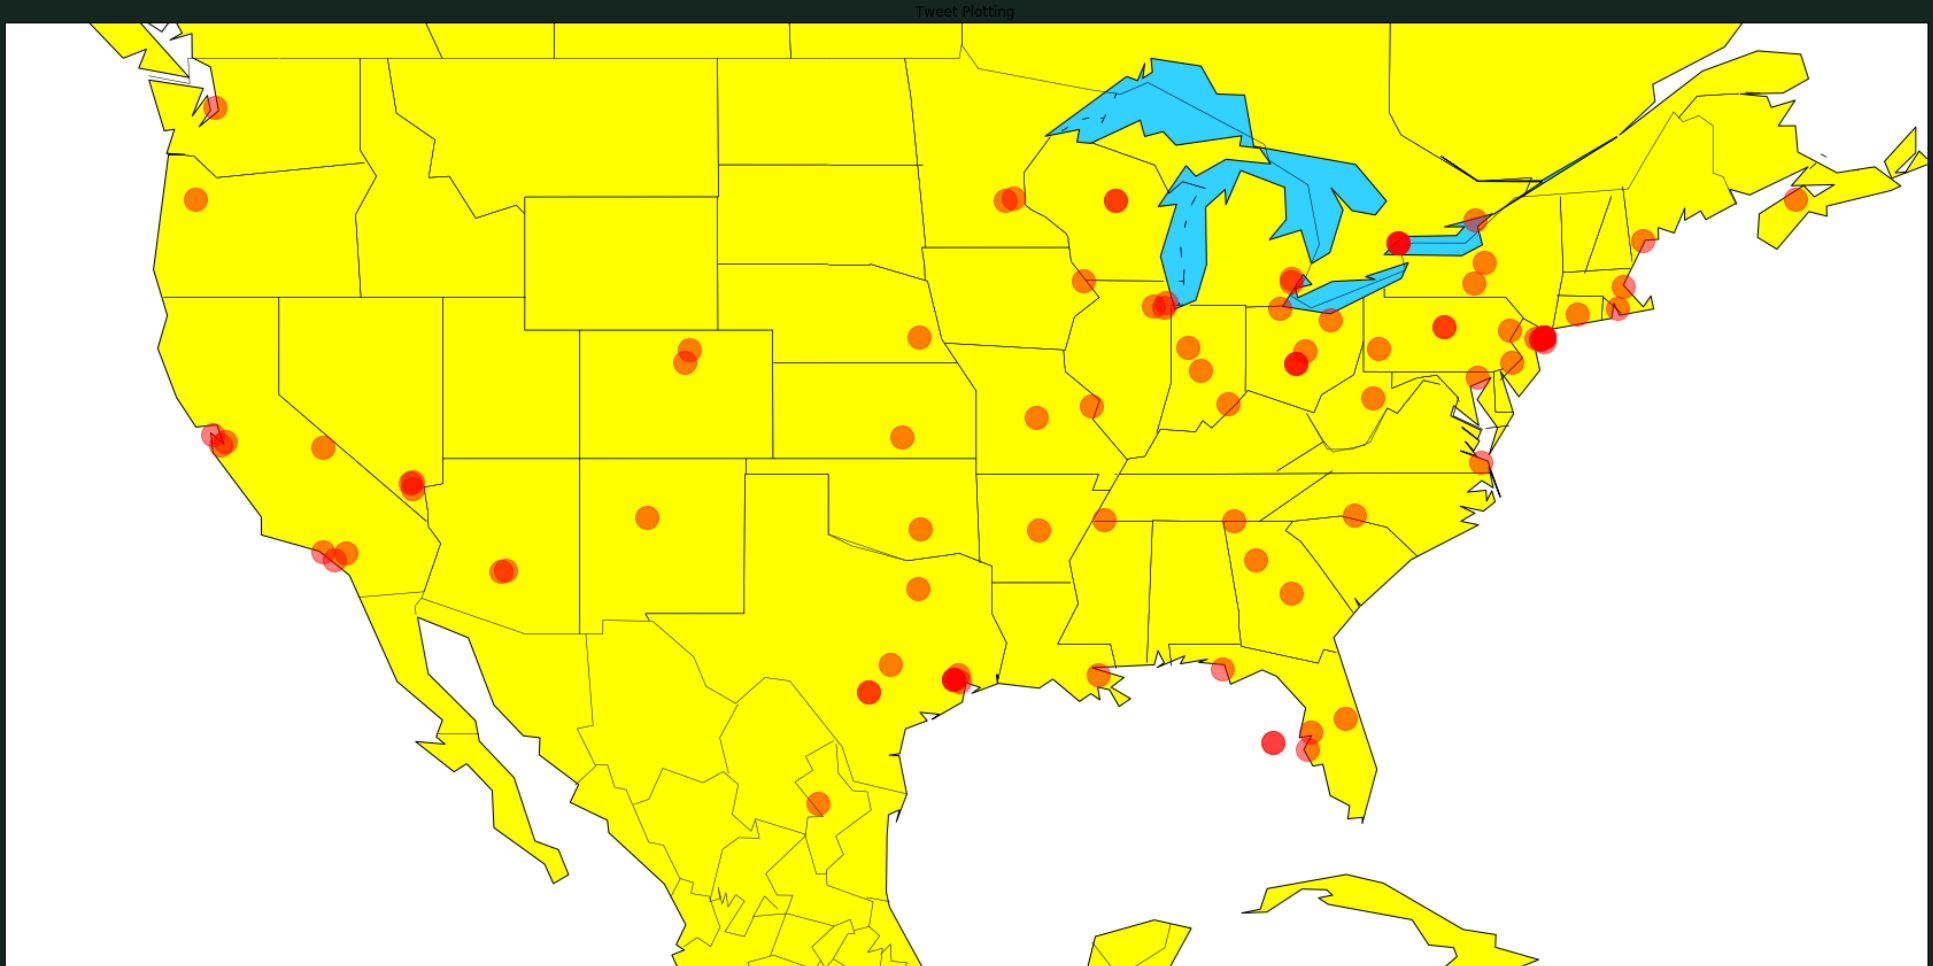
- message

- place {approx coordinates}

- retweet count

- user

Plotting Disaster locations



Best way to reach a disaster cluster

- drive a truck through the zones (road conditions, avoid danger)
- airdrop supplies (subject to getting stolen, accuracy of drop)

Clustering methods to
consider

K-means
DBScan
optical cluster

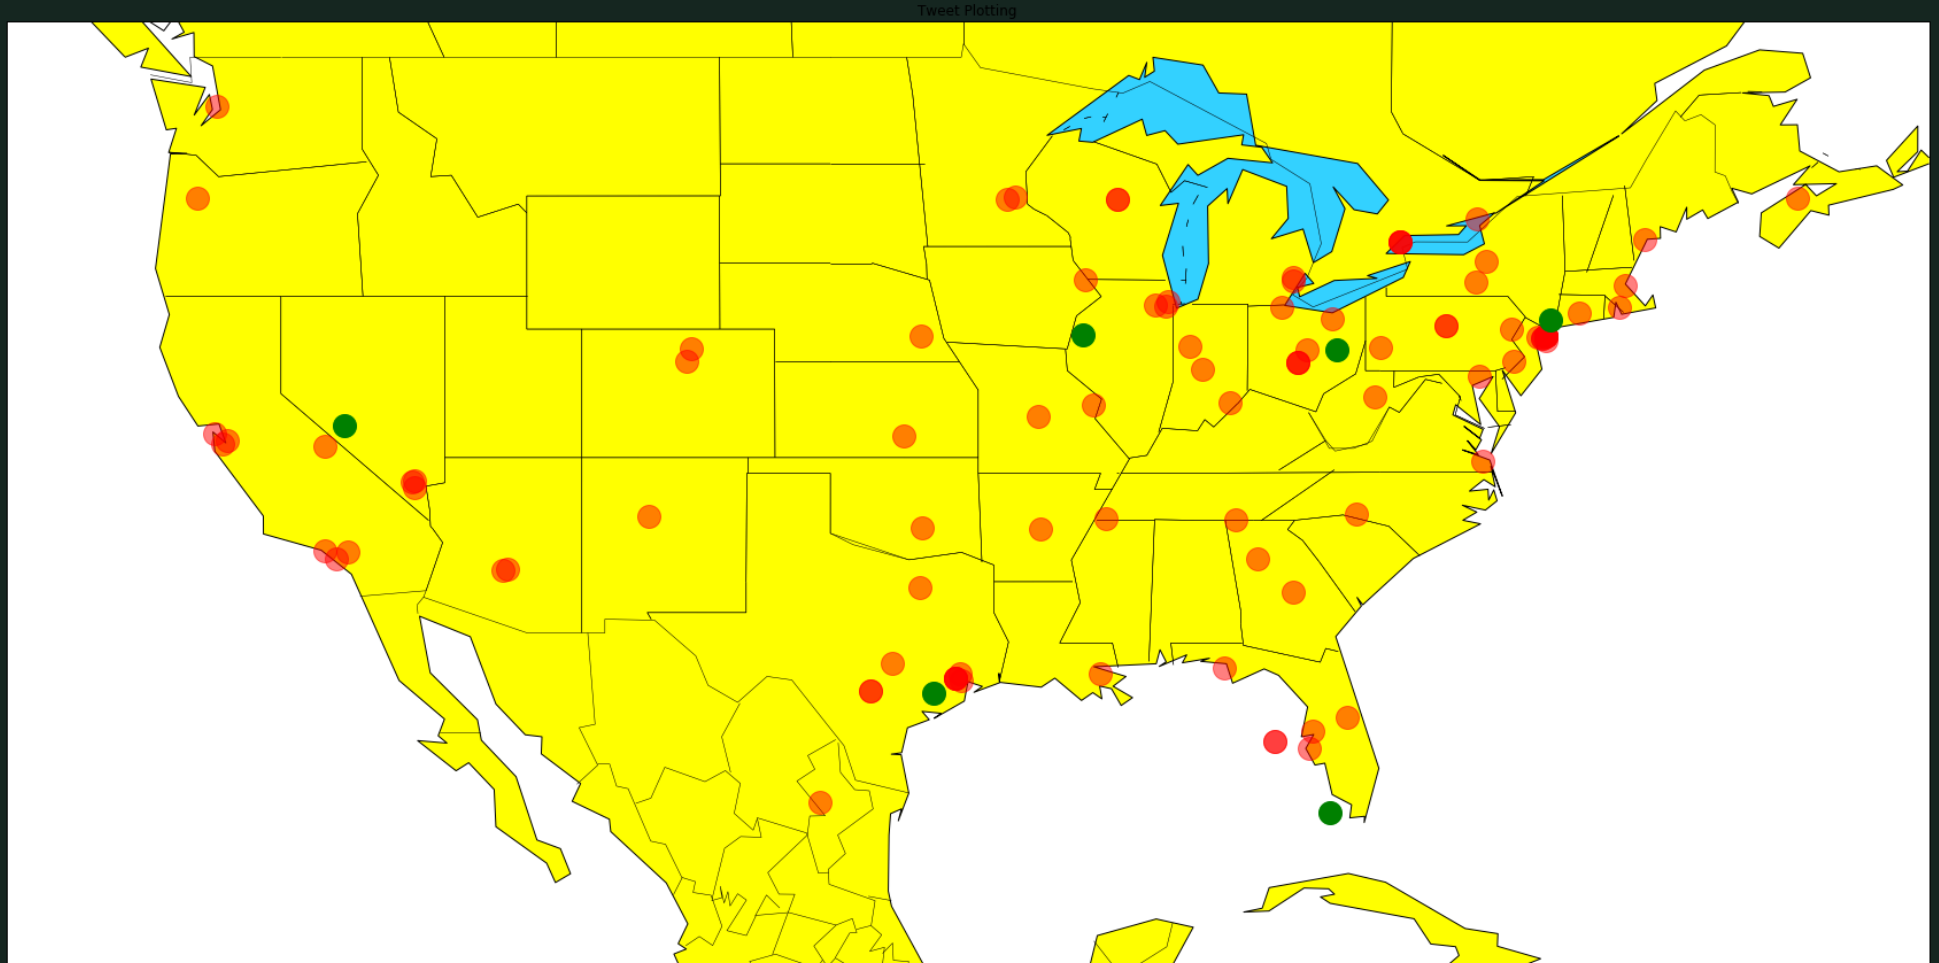
K-means wins ?

Due to the nature of the tweets , DBSCAN gave very low silhouette scores. This is due to the fact that in this scenario, we are looking to cluster points and each "rescuer" can only accomodate so many rescuees. As such, it is easier for us to approach clustering using K-means where we can set the number of clusters (rescuers) and their size.

```
In [58]: 1 start_time = time.time()
2 db = DBSCAN(eps=epsilon, min_samples=3, algorithm='auto', metric='haversine').fit(np.radians(coords))
3 cluster_labels = db.labels_
4
5 # get the number of clusters
6 num_clusters = len(set(cluster_labels))
7
8 # all done, print the outcome
9 message = 'Clustered {:,} points down to {:,} clusters, for {:.1f}% compression in {:.2f} seconds'
10 print(message.format(len(coords), num_clusters, 100*(1 - float(num_clusters) / len(coords)), time.time()-start_time)
11 print('Silhouette coefficient: {:.03f}'.format(metrics.silhouette_score(coords, cluster_labels)))
```

```
Clustered 254 points down to 11 clusters, for 95.7% compression in 0.01 seconds
Silhouette coefficient: -0.364
```

Rescue point plotting



Reasonably accurate address to begin rescue operations

```
8 print(location.address)
9
10 print((location.latitude, location.longitude))
11
12 print(location.raw)
```

Key West, Monroe County, Florida, 33040, United States of America
(24.5625566, -81.7724368)

The main rescue point for the points in
cluster 15